Partial Order Resolution of Event Logs for Process Conformance Checking

Han van der Aa[a], Henrik Leopold[b,c], Matthias Weidlich[a]

[a]*Department of Computer Science, Humboldt-Universität zu Berlin, Berlin, Germany*
[b]*Kühne Logistics University, Hamburg, Germany*
[c]*Hasso Plattner Institute, University of Potsdam, Potsdam, Germany*

Abstract

While supporting the execution of business processes, information systems record event logs. Conformance checking relies on these logs to analyze whether the recorded behavior of a process conforms to the behavior of a normative specification. A key assumption of existing conformance checking techniques, however, is that all events are associated with timestamps that allow to infer a total order of events per process instance. Unfortunately, this assumption is often violated in practice. Due to synchronization issues, manual event recordings, or data corruption, events are only partially ordered. In this paper, we put forward the problem of partial order resolution of event logs to close this gap. It refers to the construction of a probability distribution over all possible total orders of events of an instance. To cope with the uncertainty of real-world data, we present several estimators for this task, incorporating different notions of behavioral abstraction. Moreover, to reduce the runtime of conformance checking based on partial order resolution, we introduce an approximation method that comes with a bounded error in terms of accuracy. Our experiments with real-world and synthetic data reveal that our approach improves accuracy over the state-of-the-art considerably, reducing the average error by 60%.

Keywords: Process mining, Conformance checking, Data uncertainty

1. Introduction

The execution of business processes is these days supported by information systems [1]. Whether it is the handling of a purchase order in e-commerce, the tracking of an issue in complaint management, or monitoring of patient pathways in healthcare, information systems track the progress of processes in terms of event data. An event hereby denotes the execution of a specific activity (e.g., checking plausibility of a purchase order, proposing some issue resolution, creating a patient treatment plan) as part of a specific case (e.g., a purchase

order, an issue ticket, a patient) at a specific point in time [2]. A collection of such events, referred to as an event log, therefore represents the *recorded* behavior of a process.

In their efforts to improve efficiency and effectiveness of business process execution, organizations aim at standardization by means of normative process specifications [3]. That is, the activities of a process along with causal and temporal dependencies for their execution as part of a case are defined explicitly and formalized in a process model. Such a model then reflects the *desired* behavior of a process.

A considerable threat to process improvement initiatives is *non-conformance* in process execution, i.e., situations in which the recorded behavior of a process deviates from the desired one [4]. Such differences stem from the fact that information systems *support* process execution (e.g., through data storage, progress notifications, or task lists), but do not *enforce* a particular way of executing the process [5]. Rather, human interaction drives a process, giving people a certain flexibility in the execution of a particular case.

The implications of non-conformance are known to be severe. They range from reduced productivity [6] to financial penalties imposed by authorities [7]. To efficiently detect cases of non-conformance, techniques for *conformance checking* have been introduced [4, 8–10]. They strive for automatic detection of deviations of the recorded and desired process behavior, by comparing event logs with process models. They verify whether the causal dependencies for activity execution, as specified in a process model, hold true in an event log and provide diagnostic information on non-conformance.

However, a key assumption of state-of-the-art conformance checking techniques is that all events of a case are labeled with timestamps that allow to infer a *total* order [11–13]. Unfortunately, this assumption is often violated. In practice, there are various sources of uncertainty for the recorded event data, among them synchronization issues, manual recording of events, or unreliable data sensing [14]. For instance, in healthcare processes, only the day of a set of treatments may be known, but not the specific point in time [15]. Hence, events are only *partially* ordered, which renders existing conformance checking techniques inapplicable.

In this paper, we argue that it is often possible to resolve the unknown order of such events. Our idea is to use information from the entire event log to estimate the probability of each possible total order, induced by the partial order of events. This way, conformance checking is grounded in a stochastic model, incorporating the probabilities of specific order resolutions. In the presence of uncertain event data, however, diverse levels of abstraction may be needed to correlate the events of different cases for the purpose of order resolution, and it is unclear how to choose an appropriate abstraction for a given event log.

Our contributions and the structure of the paper, following background on conformance checking and uncertain event data in the next section, are summarized as follows:

2

- We introduce the problem of partial order resolution for event logs (Section 3). We formalize the problem and outline how it enables probabilistic conformance checking.

- We present various behavioral models to address partial order resolution (Section 4). These models encode different levels of abstraction of event orders, which are then used to correlate events of different cases.

- To improve the computational efficiency of conformance checking in the presence of uncertain event data, we propose a sample-based approximation method that provides statistical guarantees on obtained conformance checking results (Section 5).

- The accuracy and efficiency of our approach, as well as the approximation method are demonstrated through evaluation experiments based on real-world and synthetic data collections (Section 6). The conducted experiments reveal that our approach achieves considerably more accurate results than the state-of-the-art, reducing the average error by 60.2%.

Finally, we review our contributions in light of related work (Section 7), before concluding the paper (Section 8).

2. Background

Section 2.1 first introduces a running example and illustrates the goal of conformance checking. Section 2.2 then discusses the impact that event data uncertainty on this task.

2.1. Conformance Checking

Conformance checking analyzes deviations between the recorded and the desired behavior of a process. Recorded behavior is given as a log of events, each event carrying at least a timestamp, a reference to an activity, and an identifier of the case for which an activity was executed. Based on the latter, a log can be partitioned into *traces*: ordered, maximal sets of events, all related to the same, individual case.

The desired behavior of a process, in turn, is captured by a normative specification, i.e., a process model. It defines causal dependencies for the activities of a process, thereby inducing a set of *execution sequences*, i.e., sequences of possible activity executions that are allowed according to the process model. Conformance checking determines whether a recorded trace corresponds to an execution sequence of the process model. Put differently, it assesses whether a trace represents a *word* of the *language* of the process model.

For illustration, consider the process model depicted in Figure 1, henceforth referred to as model M. It defines the desired behavior of a healthcare process, using the Business Process Model and Notation (BPMN).

Figure 1: Process model capturing the desired behavior for a healthcare process.

That is, in each case, a patient's history should first be reviewed (activity A), before checking their insurance coverage (B) and creating a treatment plan (C). In parallel, the general practitioner (GP) of the patient is contacted (D). Depending on the result of these activities, a patient either gets a follow-up appointment (E) or is forwarded to the inpatient ward (F), before the outcome is recorded (G).

Next to this process model, consider the following traces[1]: $\pi_1 = \langle a, b, c, d, e, g \rangle$, $\pi_2 = \langle a, b, c, e, d, g \rangle$, and $\pi_3 = \langle a, d, b, f, e, g \rangle$. We observe that π_1 represents a proper execution sequence of the process model M, so we can conclude that π_1 conforms to M. By contrast, π_2 and π_3 do not. In π_2, activity E occurs *before* D. That is, a follow-up appointment was scheduled (E) without first contacting the patient's GP (D), even though the model explicitly specifies that these activities should occur in the reverse order. Trace π_3 has several different issues. First, the creation of a treatment plan (C) has been omitted, even though this represents a mandatory activity. Second, a follow-up appointment has been scheduled (E), even though the patient has also been forwarded to the inpatient ward (F). According to the model, these activities are defined to be mutually exclusive, therefore leading to another conformance issue.

Conformance checking techniques aim to automatically detect such deviations between recorded and desired process behavior. State-of-the-art techniques for this task construct *alignments* between traces and execution sequences of a model to detect deviations [4, 12, 13, 16]. An alignment is a sequence of steps, each step comprising a pair of an event and an activity, or a *skip* symbol $\bot$, if an event or activity is without counterpart. For instance, for the non-conforming trace π_3, an alignment with two such skip steps may be constructed with the execution sequence $\langle A, D, B, C, F, G \rangle$ of the model:

Trace π_3	a	d	b	$\bot$	f	e	g
Execution sequence	A	D	B	C	F	$\bot$	G

Assigning costs to skip steps, a cost-optimal alignment (not necessarily unique) is constructed for a trace in relation to all execution sequences of a model [12]. An optimal alignment then answers not only the question

---

[1]We use lowercase letters to refer to events corresponding to a certain activity, e.g., e describes an occurrence of activity E.

whether there are deviations between a trace and the execution sequences of a model, but also enables quantification of non-conformance by aggregating the costs of skip steps. In addition, considering a log as a whole, events and activities that are frequently part of skip steps highlight hotspots of non-conformance in process execution.

2.2. Event Data Uncertainty

While various conformance checking techniques have been presented in recent years, they all assume and require event data to have been accurately recorded. Yet, in practice, event logs are subject to diverse quality issues [14, 17, 18] along all the dimensions known to assess data quality in general [19], e.g., accuracy, timeliness, precision, completeness, and reliability.

In this work, we focus on quality issues that relate to temporal aspects of events, which are highly relevant to the conformance checking task. In particular, state-of-the-art conformance checking relies on discrete events that are assigned a precise timestamp. Hence, the commonly adopted notion of a trace requires all events of a single case to be *totally* ordered by their timestamps [12, 16]. However, this assumption is often violated in practice, for reasons such as:

1. *Lack of synchronization.* Event logs integrate data from various information systems. Tracing the execution in such distributed systems has to cope with unsynchronized clocks [20], making timestamps partially incomparable. Also, the logical order of events induced by timestamps may be inconsistent with the order of their recording [21].

2. *Manual recording.* The execution of activities is not always directly observed by information systems. Rather, people involved in process execution have to record them manually. Such manual recordings are subject to inaccuracies. For instance, it has been observed in the healthcare domain that personnel records their work solely at the end of a shift [22], rendering it impossible to determine a precise order of executed activities.

3. *Data sensing.* Event logs may be derived from sensed data as recorded by real-time locating systems (RTLS). Then, the construction of discrete events from raw signals is inherently uncertain and grounded in probabilistic inference [23]. For instance, deriving treatment events in a hospital based on RTLS positions of patients and staff members does not yield fully accurate traces [24].

The above phenomena have in common that they result in imprecise event timestamps. In conformance checking, this leads to the particular problem of *order uncertainty*: The exact order in which events of a

5

Table 1: Events recorded for a single case.

Event ID	Activity	Timestamp
a	Review patient history [A]	13:00
b	Check insurance coverage [B]	14:00
c	Create treatment plan [C]	14:00
f	Forward to inpatient ward [F]	15:00
d	Contact patient's GP [D]	15:00
g	Record outcome [G]	16:00

trace have occurred is not known. Consider, for instance, the events shown in Table 1, recorded for a single case of the aforementioned process. The events may have been captured manually, so that the timestamps only indicate the rough hour in which activities have been executed. Since events b and c carry the same timestamp, it is unclear whether the patient's insurance was checked (B) *before* or *after* creating a treatment plan (C). Since the type of insurance may influence the treatment plan, the model in Figure 1 defines an explicit execution order for both activities. Yet, due to order uncertainty, we cannot establish whether the process was indeed executed as specified.

A result of order uncertainty, whether caused by a lack of synchronization, manual recording, or data sensing issues, is that the events of a trace are only *partially* ordered. Such a partial order is visualized in Figure 2 for the events of the example case. This partial order induces four totally ordered sequences of events, denoted as π_4 to π_7 in the figure. Such a situation is highly problematic in conformance checking, because of the implied ambiguity. For the example in Table 1, only one of the totally ordered event sequences, i.e., π_4 conforms to model M, whereas different kinds of deviations are detected for the remaining three. Hence, one cannot conclude if the case was executed as specified by the model at all and, if it was not, which conformance violations actually occurred.

Without further insights into the execution of a case, two approaches may be followed to resolve order uncertainty. First, one may consider all induced totally ordered traces to be equally likely, so that the number of such traces that conform to the model provides a conformance measure. Second, order uncertainty may be neglected, verifying whether *one* of the induced total orders is conforming [22]. As we will demonstrate empirically, both approaches introduce a severe bias in conformance checking. In this work, we therefore strive for a fine-granular assessment of each induced total event order.


$$\pi_4 = \langle a, b, c, d, f, g \rangle$$
$$\pi_5 = \langle a, c, b, d, f, g \rangle$$
$$\pi_6 = \langle a, b, c, f, d, g \rangle$$
$$\pi_7 = \langle a, c, b, f, d, g \rangle$$

Figure 2: Partial order and trace resolutions resulting from the order uncertainty in the case of Table 1.

3. Problem Statement

We first give preliminaries in terms of a formal model (Section 3.1), before defining the problem addressed in this paper, partial-order resolution (Section 3.2).

3.1. Preliminaries

In this section, we introduce our formal model to capture normative and recorded process behavior.

Normative process behavior. A process model defines the execution dependencies between the activities of a process, establishing the normative or desired behavior. For our purposes, it is sufficient to abstract from specific process modeling languages (e.g., BPMN or Petri nets) and focus on the behavior defined by a model. Process models capture relations that exist among a collection of *activities*. We denote the universe of such activities as $\mathcal{A}$. Then, a process model defines a set of execution sequences, $M \subseteq \mathcal{A}^*$, that capture sequences of activity executions that lead the process to its final state. For instance, the model in Figure 1 defines a total of six allowed execution sequences, including $\langle A, B, C, D, E, G \rangle$, $\langle A, B, C, D, E, F \rangle$, as well as variations in which activity D occurs before or after B.

Recorded process behavior. The executions of activities of a process are recorded as events. If these activity executions happen within the context of a single case, the respective events are part of the same *trace*. While most models define traces as sequences of events, we adopt a model that explicitly captures order uncertainty by allowing multiple events, even if belonging to the same trace, to be assigned the same timestamp. Hence, the events of a trace are only *partially* ordered. We capture this by modeling traces as sequences of *sets* of events, where each set contains events with identical timestamps. Captured as follows:

Definition 1 (Traces). *A trace is a sequence of disjoint sets of events, $\sigma = \langle E_1, \ldots, E_n \rangle$, with $E_\sigma = \bigcup_{1 \leq i \leq n} E_i$ as the set of all events of σ.*

For a trace $\sigma = \langle E_1, \ldots, E_n \rangle$, we refer to E_i, $1 \leq i \leq n$, as an *event set* of σ. This event set is *uncertain*, if $|E_i| > 1$. Accordingly, a trace that does not contain uncertain event sets is *certain*, otherwise it is *uncertain*. Intuitively, in a certain trace, a total order of events is established by the events' timestamps. For an uncertain trace, events within an uncertain event set are not ordered.

An event log is a set of traces, capturing the events as they have been recorded during process execution. Moreover, for each event, we capture the activity for which the execution is represented by this event. The latter establishes a link between an event log and the activities of a process model.

Definition 2 (Event Log). *An event log is a tuple $L = (\Sigma, \lambda)$, where Σ is a set of traces and $\lambda : \bigcup_{\sigma \in \Sigma} E_\sigma \to \mathcal{A}$ assigns activities to all events of all traces.*

As a short-hand notation, we write a trace of a log not only as a sequence of sets of events, but also as a sequence of sets of activities. That is, for a trace $\langle \{e_1, e_2\}, \{e_3\} \rangle$ with $\lambda(e_1) = x$, $\lambda(e_2) = y$, and $\lambda(e_3) = z$, we also write $\langle \{x, y\}, \{z\} \rangle$. According to this model, the case from Table 1 is captured by the trace $\sigma_1 = \langle \{a\}, \{b, c\}, \{d, f\}, \{g\} \rangle$.

3.2. The Partial Order Resolution Problem

To assess the conformance of an event log with a model, the order uncertainty of its traces needs to be handled. Yet, there may be several ways to resolve this uncertainty as the events of each uncertain event set may be ordered differently. We capture such different orders by means of *possible resolutions* of event sets and, based thereon, of a trace.

Definition 3 (Possible Resolutions). *Given a trace $\sigma = \langle E_1, \ldots, E_n \rangle$, we define* possible resolutions *for:*
- *an event set E_i, as any total order over its events, i.e., $\Phi(E_i) = \{ \langle e_1, \ldots, e_{|E_i|} \rangle \mid \forall\ 1 \leq j, k \leq |E_i| : e_j \in E_i \wedge e_j = e_k \Rightarrow j = k \}$;*
- *the trace σ, as any total order over events of its event sets, i.e., $\Phi(\sigma) = \{ \langle e_1^1, \ldots, e_1^{m_1}, \ldots, e_n^1, \ldots e_n^{m_n} \rangle \mid \forall\ 1 \leq i \leq n : \langle e_i^1, \ldots e_i^{m_i} \rangle \in \Phi(E_i) \}$.*

In the context of an event log $L = (\Sigma, \lambda)$, we lift the short-hand notation for traces based on the assigned activities to resolutions. Then, for our example trace $\sigma_1 = \langle \{a\}, \{b, c\}, \{d, f, \}, \{g\} \rangle$, possible resolutions would be $\langle a, c, b, f, d, g \rangle$ or $\langle a, b, c, d, f, g \rangle$, but neither $\langle a, b, f, d, g \rangle$ (event c is missing) nor $\langle a, c, f, b, d, g \rangle$ (events from different event sets are interleaved).

Although the events originally occurred in a total order, there is no way to recover this original order when it is obscured due to the aforementioned reasons for order uncertainty. However, we argue that even without identifying a single resolution, valuable insights on the conformance of a trace may be obtained. This can be achieved by assessing which resolutions of a trace conform and which do not conform to a process model. As illustrated above, it is crucial here to avoid basing conformance assessments purely on the *number* of possible resolutions that conform to its associated process model: a single conforming resolution may be more likely to have occurred than multiple non-conforming resolutions combined.

We therefore resort to a probabilistic model that defines a probability distribution over a trace's possible resolutions. Then, conformance checking can be grounded in the cumulative probabilities of the resolutions that conform to a model, or show particular deviations, respectively. Following this line, a crucial problem addressed in this paper is how to assign probabilities to the possible resolutions of a partial order, which can be phrased as follows:

Problem 1 (Partial Order Resolution). *Let $L = (\Sigma, \lambda)$ be an event log. The* partial order resolution *problem is to derive, for each trace $\sigma \in \Sigma$ and each possible resolution $\varphi \in \Phi(\sigma)$, the probability $P(\varphi)$ of φ representing the order of event generation for the respective case.*

Based on the probabilities $P(\varphi)$ of each resolution $\varphi \in \Phi(\sigma)$, the probabilistic conformance of a trace σ with respect to a model M can be assessed. Let $conf(\varphi, M)$ be a function that quantifies the conformance of a resolution to model M. Exemplary functions are a binary function $conf_{bin} : \varphi \times M \rightarrow \{0, 1\}$, indicating a conforming resolution with 1 and a non-conforming with 0, or a function based on trace fitness [25], yielding a range from 0 to 1, i.e., $conf_{fit} : \varphi \times M \rightarrow [0, 1]$. Based on such a function, the weighted conformance of a trace is denoted as follows:

$$P_{conf}(\sigma, M) = \sum_{\varphi \in \Phi(\sigma)} P(\varphi) \times conf(\varphi, M) \qquad (1)$$

Similarly, more fine-granular feedback based on non-conformance may be given by analyzing alignments obtained per possible resolution of a trace. For instance, the probabilities assigned to resolutions can be incorporated as weighting factors in the aggregation of the non-conformance measured per possible resolution. This manifests itself in the form of the accumulative probabilities associated with the skip steps ($\perp$) in an alignment, as shown in Section 2.1. In this way, probabilistic conformance checking can be used to identify hotspots of non-conformance in processes.

4. Partial Order Resolution

To estimate the probability of a partial order resolution, we follow the idea that the context of a trace provided by the event log is beneficial. A business process is structured through the causal dependencies for the execution of activities. These dependencies are manifested in the traces in terms of behavioral regularities. Assuming that order uncertainty occurs independently of the execution of the process, behavioral regularities among the traces may be exploited for partial order resolution. The probability of a specific resolution for a given trace may be assessed based on order information derived from similar traces contained in the event log. Intuitively, if one possible resolution denotes an order of activity executions that is frequently observed for other certain traces, this resolution is expected to be more likely than another resolution that denotes a rare sequence of activity executions. It is important to note that model characteristics cannot be leveraged to determine the likelihood of a resolution since this would introduce a bias towards conforming resolutions.

Using traces for partial order resolution requires a careful selection of the abstraction level based on which

9

Table 2: Proposed behavioral models.

Model	Illustration	Basis
Trace equivalence	*full trace* < a, b, c, d, e >	Equal, certain traces
N-gram	*N-gram* < a, b, c, d, e >	Equal sub-sequences of length N
Weak order	*weak order* < a, b, c, d, e >	Indirectly follows relation of events

traces are compared. In practice, event logs contain traces encoding a large number of different sequences of activity executions. Reasons for that are concurrent execution of activities, which leads to an exponential blow-up of the number of execution sequences, as well as the presence of noise, such as incorrectly recorded events. For a possible resolution of a trace, it may therefore be impossible to observe the exact same sequence of activity executions in another trace, unaffected by order uncertainty.

We cope with this issue by defining behavioral models that realize different levels of abstraction in the comparison of traces. We propose (i) the trace equivalence model, (ii) the N-gram model, and (iii) the weak order model, see Table 2. The models differ in the notion of behavioral regularity that is used for the partial order resolution.

4.1. Trace Equivalence Model

This model estimates the probability of a resolution by exploring how often the respective sequence of activity executions is observed in the event log, in traces without order uncertainty. To this end, we first clarify that two resolutions shall be considered to be equivalent, if they represent the same sequences of executed activities.

Let $L = (\Sigma, \lambda)$ be an event log and $\sigma, \sigma' \in \Sigma$ two traces of the same length, i.e., $|E_\sigma| = |E_{\sigma'}|$. Let $\varphi = \langle e_1, \ldots, e_n \rangle \in \Phi(\sigma)$ and $\varphi' = \langle e_1', \ldots, e_n' \rangle \in \Phi(\sigma_c)$ two of their resolutions. Then, the resolutions are equivalent, denoted by $\varphi \equiv \varphi'$, if and only if $\lambda(e_i) = \lambda(e_i')$ for $1 \leq i \leq n$.

We define $\Sigma_{certain} = \{\sigma \in \Sigma \mid |\Phi(\sigma)| = 1\}$ as the set of all certain traces, i.e., all traces that do not have uncertainty and, thus, only a single resolution. Then, we quantify the probability associated with a resolution $\varphi \in \Phi(\sigma)$ of a trace σ as the fraction of certain traces for which the resolutions are equivalent:

$$P_{trace}(\varphi) = \frac{|\{\sigma \in \Sigma_{certain} \mid \exists \varphi' \in \Phi(\sigma) : \varphi' \equiv \varphi\}|}{|\Sigma_{certain}|} \tag{2}$$

10

The above model enables a direct assessment of the probability of a resolution. Yet, it may have limited applicability: (i) It only considers certain traces, and there may only be a small number of those in a log; (ii) none of the certain traces may show an equivalent resolution, as it requires the *entire* sequence of activity executions to be the same.

4.2. N-Gram Model

As a second approach, we introduce a behavioral model based on *N-gram approximation.* It takes up the idea of sequence approximations as they are employed in a broad range of applications, such as prediction [26] and speech recognition [27]. Specifically, N-gram approximation enables us to define a more abstract notion of behavioral regularities that determines which traces shall be considered when computing the probability of a resolution.

Given a resolution of a trace, this model first estimates the probability of the individual events of the resolution occurring at their specific position. Here, up to $N-1$ events preceding the respective event are considered and their probability of being followed by the event in question is determined. The latter is based on all traces of the log that comprise sub-sequences of the same activity executions without ordering uncertainty.

For instance, for a resolution $\varphi_1 = \langle a, b, c, d, f, g \rangle$, we determine how likely the event f occurs at the fifth position of the resolution. Setting $N = 4$, we would then explore how likely the sequence of activity executions $\langle b, c, d \rangle$ is followed by the execution of f. This estimation is based on all traces of the log that comprise events of the same sequence of activity executions without order uncertainty.

To formalize this idea, we first define a predicate *certain.* Given a log $L = (\Sigma, \lambda)$, this predicates holds for a sequence of activities $\langle a_1, \ldots, a_m \rangle$, $a_i \in \mathcal{A}$ for $1 \leq i \leq m$ and a trace $\sigma = \langle E_1, \ldots, E_n \rangle \in \Sigma$, $m \leq n$, if σ contains events for the respective activity executions without order uncertainty:

$$certain(\langle a_1, \ldots, a_m \rangle, \sigma = \langle E_1, \ldots, E_n \rangle) \Leftrightarrow$$
$$\exists\, i \in \{0, \ldots, n - m\}, \forall\, j \in \{1, \ldots, m\} : E_{i+j} = \{e_{i+j}\} \wedge \lambda(e_{i+j}) = a_j. \tag{3}$$

Using this predicate, we define the probability of events related to activity a to follow events denoting the execution of some activities $\langle a_1, \ldots, a_m \rangle$. This definition is based on the number of times the two respective sequences, with and without a, are observed in the traces of the event log:

$$P(a \mid \langle a_1, \ldots, a_m \rangle) = \frac{|\{\sigma \in \Sigma \mid certain(\langle a_1, \ldots, a_m, a \rangle, \sigma)\}|}{|\{\sigma \in \Sigma \mid certain(\langle a_1, \ldots, a_m \rangle, \sigma)\}|}. \tag{4}$$

For illustration, we return to the example of estimating the probability of events related to f to be preceded by those representing the execution of activities $\langle b, c, d \rangle$. That is, we divide the number of occurrences of $\langle b, c, d, f \rangle$ by the number of occurrences of $\langle b, c, d \rangle$, while considering only traces that do not show order uncertainty for the respective events. Based thereon, the probability of a resolution is derived by aggregating the N-gram-based probabilities of all its events:

$$P_{N\text{-}gram}(\varphi = \langle e_1, \ldots e_n \rangle) = \prod_{k=2}^{n} P(\lambda(e_k) \mid \langle \lambda(e_{\max(1, k-N+1)}), \ldots, \lambda(e_{k-1}) \rangle) \tag{5}$$

The above approach may be adapted to explicitly consider the first events of traces in the assessment. Technically, an artificial event is added to the beginning of all traces, so that it will be part of the respective N-gram definitions. For instance, the example trace $\sigma_1 = \langle \{a\}, \{b, c\}, \{d, f, \}, \{g\} \rangle$ would be changed to $\sigma_1' = \langle \{\circ\}, \{a\}, \{b, c\}, \{d, f, \}, \{g\} \rangle$ with $\circ$ denoting the start of the trace. Then, the estimation of the probability for the resolution $\varphi_1' = \langle \circ, a, b, c, d, f, g \rangle$ would, using $N = 4$, be based on an assessment of the probability that c is preceded by $\langle \circ, a, b \rangle$. This explicitly considers solely traces that start with events that denote executions of a and b.

Compared to the trace equivalence model, the N-gram model is more abstract. This makes the model more generally applicable, as it requires only the presence of traces that show equivalent sub-sequences of activity executions, not involving order uncertainty, instead of equivalent traces. The parameter N provides further flexibility. Higher values of N lead to longer sub-sequences being considered, which induces a stricter notion of behavioral regularities to be exploited in partial order resolution. Lower values, in turn, decrease this strictness, thereby increasing the amount of evidence on which the resolution is based.

Yet, the N-gram model assumes that behavioral regularities materialize in the form of consecutive activity executions. Even if $N = 2$, only events that (certainly) follow upon each other directly in a trace are considered in the assessment.

4.3. Weak Order Model

To obtain an even more abstract model, we drop the assumption that behavioral regularities relate solely to consecutive executions of activities. Rather, indirect order dependencies, referred to as a *weak order*, among the activity executions, and thus events, are exploited.

To illustrate this idea, consider two resolutions, $\varphi_1 = \langle a, b, c, d, f, g \rangle$ and $\varphi_2 = \langle a, c, b, d, f, g \rangle$, of our

example trace σ_1. To estimate their probabilities, under a weak order model, we determine the fraction of traces in which an event related to activity b occurs at some point before c, or vice versa, to obtain evidence about the most likely order. Assume that the event log also contains an uncertain trace $\sigma_2 = \langle\{a,b\},\{d\},\{c\},\{f,g\}\rangle$. In this trace, b and c are not part of consecutive event sets, and b is even part of an uncertain event set. Still, this trace provides evidence that activity b is executed before c, which supports resolution φ_1 of trace σ_1. This model also enables us to incorporate information from consecutive uncertain event sets, as in $\sigma_1 = \langle\{a\},\{b,c\},\{d,f\},\{g\}\rangle$. Due to order uncertainty, it is unclear whether events c and f directly followed each other. However, c definitely occurred earlier than f. Such information would not be taken into account by the N-gram model.

Formally, let $L = (\Sigma, \lambda)$ be an event log and $a, a' \in \mathcal{A}$ two activities. We define a predicate *order* to capture whether a trace comprises events representing the executions of these activities in weak order:

$$order(a, a', \sigma = \langle E_1, \ldots, E_n\rangle) \Leftrightarrow \exists\, i, j \in \{1, \ldots, n\}, i < j : e_i \in E_i \wedge e_j \in E_j \wedge \lambda(e_i) = a \wedge \lambda(e_j) = a'. \quad (6)$$

This predicate enables us to estimate the probability of having events related to specific activity executions in weak order. More specifically, we determine the ratio of traces that contain the respective events:

$$P(a, a') = \frac{|\{\sigma \in \Sigma \mid order(a, a', \sigma)\}|}{|\{\sigma \in \Sigma \mid \exists\, e, e' \in E_\sigma : \lambda(e) = a \wedge \lambda(e') = a'\}|}. \quad (7)$$

Based thereon, the probability of a resolution is defined by aggregating the probabilities of all pairs of events to occur in the particular order:

$$P_{WO}(\varphi = \langle e_1, \ldots e_n\rangle) = \prod_{\substack{1 \leq i < n \\ i < j \leq n}} P(\lambda(e_i), \lambda(e_j)). \quad (8)$$

Compared to the other two models, the weak order model employs the most abstract notion of behavioral regularity when resolving partial orders. Consequently, the computation of the probability of a particular resolution can be expected to exploit information from many traces of the event log.

5. Result Approximation

A key issue hindering the applicability of conformance checking techniques in industry is their computational complexity, since state-of-the-art algorithms suffer from an exponential runtime complexity in the length of the trace [4]. In the context of this paper, this is particularly problematic: Order uncertainty exponentially

increases the number of conformance checks that are required per trace. To increase the applicability of conformance checking under order uncertainty, this section therefore proposes an approximation method incorporating statistical guarantees.

This section discusses the calculation of expected conformance values (Section 5.1) and associated confidence intervals (Section 5.2), before describing the algorithmic computation method utilized in our approach (Section 5.3).

5.1. Expected Conformance

To reduce the computational complexity of the conformance checking task, we propose an approximation method that provides statistical guarantees about the conformance results $P_{conf}(\sigma, M)$ obtained for a trace σ. The method provides a confidence interval for the conformance results, which is computed based on conformance checks performed for a sample of its possible resolutions $\Phi(\sigma)$. We use $\bar{\Phi} \subseteq \Phi(\sigma)$ to denote this sample and $\bar{p} = \sum_{\varphi \in \bar{\Phi}} P(\varphi)$ to denote the cumulative probability of the resolutions in the sample. Then, we define the expected conformance for a trace σ to a process model M as follows:

$$E(P_{conf}(\sigma, M)) = \sum_{\varphi \in \bar{\Phi}} P(\varphi) \times conf(\varphi, M) + (1 - \bar{p}) \times \mu_{conf} \tag{9}$$

Equation 9 consists of two components: (i) the known, weighted conformance of the sampled resolutions, i.e., $\sum_{\varphi \in \bar{\Phi}} P(\varphi) \times conf(\varphi, M)$, and (ii) an estimated part, $(1 - \bar{p}) \times \mu_{conf}$. This estimated part receives a weight of $1 - \bar{p}$, i.e., the cumulative probability of the traces not included in the sample. The estimate itself, i.e., μ_{conf}, reflects the expected conformance of a previously unseen resolution. This value is obtained by fitting a statistical distribution over the conformance values obtained for the sample $\bar{\Phi}$. Consider the conformance functions discussed in Section 3.2: For a binary function $conf_{bin}$ with range $\{0, 1\}$, the results of a sample represent a *Binomial distribution*. For a more fine-granular function $conf_{fit}$ based on trace fitness with range $[0, 1]$, the resulting distribution can be characterized using a *normal distribution* for a sufficiently large sample (e.g., size over 20) following the central limit theorem [28].

For illustration, consider a sample $\bar{\Phi}$ that contains 30 resolutions, 21 conforming and 9 non-conforming. Then, the estimated, binary conformance of unseen resolutions is given as $\mu_{conf} = 0.70$. With a cumulative probability of $\bar{p} = 0.80$, the estimated component of Equation 9 equals $(1 - 0.80) \times 0.70 = 0.14$. Note that determining the expected conformance μ_{conf} is independent of the probabilities assigned by a behavioral model. Therefore, due to differences among the probabilities associated with the resolutions in $\bar{\Phi}$, it does not necessarily hold that $E(P_{conf}(\sigma, M)) = \mu_{conf}$. For instance, for the given example, we may have

$\sum_{\varphi \in \bar{\Phi}} P(\varphi) \times conf(\varphi, M) = 0.60$, yielding an estimated overall conformance of $0.60 + 0.14 = 0.74$, which is considerably higher than μ_{conf}.

5.2. Confidence Intervals

Based on statistical distributions established for expected conformance values, we further derive statistical bounds in the form of a confidence interval for the estimation of $P_{conf}(\sigma, M)$. Recognizing that one part of the conformance checking results is known, whereas the other requires estimation, we define the confidence interval as follows:

$$CI_\alpha = E(P_{conf}(\sigma, M)) \pm (1 - \bar{p}) \times m_\alpha \tag{10}$$

Equation 10 consists of two components: (i) the expected conformance $E(P_{conf}(\sigma, M))$, and (ii) a margin, $\pm(1 - \bar{p}) \times m_\alpha$, which determines the size of the interval. Here, m_α denotes the *margin of error* of the distribution for a significance level α. The margin of error for a Binomial distribution obtained over a set of binary conformance assessments, m_α is computed using, e.g., the *Wilson score interval* [29]. For a normal distribution, the margin of error is based on the standard error [30].

It is important to note that the width of a confidence interval established using Equation 10 decreases if the cumulative probability of the sample ($\bar{p}$) is higher. This property naturally follows, because the margin of error m_α is only applicable to the estimated component of a conformance assessment, which has the weight $1 - \bar{p}$. We utilize this property in the approximation method described next.

5.3. Computation Method

Our proposed method for efficient conformance approximation is presented in Figure 1. The approximation is an iterative procedure that incorporates the conformance of newly sampled resolutions until a sufficiently accurate conformance value is established.

Loop. Each iteration starts by selecting a resolution ϕ, which has a maximal likelihood from those in $\Phi(\sigma) \setminus \bar{\Phi}$ (line 7). This selection is guided by one of the behavioral models introduced in Section 4, as configured by the input parameter B. Next, the algorithm computes the conformance result for ϕ and adds it to R, the bag of conformance results (line 9). Then, statistical distribution is fit to the results sample (line 10). Based on this distribution, the estimated result is computed using Equation 9 (line 11), before the margin of error is determined (line 12) and the accumulated probability of all sampled resolutions is updated (line 13).

Stop condition. The iterative procedure is repeated until the margin of error leads to results that are below a user-specified precision threshold δ. In particular, the method repeatedly samples additional resolutions until the ratio of the margin of error m_α and the estimated conformance E, weighted by the complement of the accumulated probability of sampled resolutions, is below δ (line 14). The stop condition is based on the ratio, rather than the absolute margin of error, given that, for instance, a margin of 0.05 has a considerably greater impact when $E = 0.2$ than compared to $E = 0.8$.

By employing this approximation method, we obtain results that satisfy a desired significance value α and precision level δ. This means that, when possible, the method requires only a relatively low number of resolutions, whereas for traces with a higher variability among its resolutions, the method will use a greater sample to ensure result accuracy.

Algorithm 1 Statistical conformance approximation method

1: **input** σ, a trace; M, a process model; B, a behavioral model; $conf$, a conformance
2:        function; α, a significance level; δ, a desired accuracy threshold.
3: $\bar{\Phi} \leftarrow \emptyset$ ▷ The set of sampled resolutions
4: $R \leftarrow [\,]$ ▷ Bag of conformance results
5: $\bar{p} \leftarrow 0$ ▷ Accumulated probability of sampled resolutions
6: **repeat**
7:        $\varphi \leftarrow select(\Phi(\sigma) \setminus \bar{\Phi}, B)$ ▷ Sample a new resolution
8:        $\bar{\Phi} \leftarrow \bar{\Phi} \cup \{\varphi\}$ ▷ Add resolution to sample
9:        $R \leftarrow R \uplus [conf(\varphi, M)]$ ▷ Compute and add conf. result
10:        $\mathcal{D} \leftarrow fit(R)$ ▷ Fit distribution on the results
11:        $E \leftarrow estimateResult(R, \mathcal{D})$ ▷ Use Equation 9
12:        $m_\alpha \leftarrow computeMargin(\mathcal{D}, \alpha)$ ▷ Margin of error
13:        $\bar{p} \leftarrow \bar{p} + P_B(\phi)$ ▷ Increase accumulated probability
14: **until** $(1 - \bar{p}) \times m_\alpha / E \leq \delta$ ▷ Check threshold
15: **return** $E \pm m_\alpha$ ▷ Return confidence interval

---

6. Evaluation

This section describes evaluation experiments in which we assess the accuracy of the proposed behavioral models for conformance checking. We achieve this by comparing the conformance checking results obtained for uncertain traces to the conformance results that would have been obtained without order uncertainty. Furthermore, we also determine which behavioral model should be used in which situation.

Both the datasets and the implementation used to conduct these evaluation experiments are publicly available.[2]

---

[2]https://github.com/hanvanderaa/uncertainconfchecking

We conducted our evaluation based on both real-world and synthetic data collections.

Real-world collection. We used three, publicly available, real-world events logs, detailed in Table 3. As shown in the table, the three logs differ considerably, for instance in terms of their size (13,087 for BPI-12 to 150,370 traces for the traffic fines log) and trace length (averages from 3.7 to 20.0 events per trace). These logs also demonstrate the prevalence of coarse-grained timestamps in real-world settings, since all logs contain a considerable amount of uncertain traces, ranging up to 97.2% of the traces for the BPI-14 log. To obtain a process model that serves as a basis for conformance checking, we use the inductive miner [31], a state-of-the-art process discovery technique, with the default parameter setting (i.e., a noise filtering threshold of 80%).

Table 3: Characteristics of the real-world collection.

Characteristic	BPI-12 [32]	BPI-14 [33]	Traffic fines [34]
Places	22	39	19
Transitions	46	64	24
Traces	13,087	46,616	150,370
Variants	4,366	31,725	231
Trace length (avg.)	20.0	7.3	3.7
Uncertain traces	5,006 (38.3%)	45,334 (97.2%)	9,166 (6.1%)
Resolutions (avg.)	3.3	21.2	95.2
Resolutions (max.)	12	768	960

Synthetic collection. We have generated a collection of 500 synthetic models that allows us to assess the impact of model characteristics such as loops and arbitrary skips on the performance of our approach, as well as to control the level of non-conformance in event logs. We employed the state-of-the-art process model generation technique from [35] due to its ability to also generate non-structured models, e.g., models that include non-local decisions. To generate models, we employed the default parameters set by the technique's developers.

For each of these models, we used a stochastic simulation plug-in of the ProM 6 framework [36] to generate an event log consisting of 1000 traces, using the default plug-in settings (i.e., uniform probabilities for all choices and exponential inter-arrival and execution times). The main characteristics of these models and their logs are detailed in Table 4.

To introduce non-conformance, we insert noise using the same simulation plug-in, which randomly inserts, swaps, and removes events in a specified fraction of the traces. For each model, we created logs with four different noise levels, by inserting noise into 25%, 50%, 75%, and 100% of the traces. This results in four sets of logs with, on average 353.8, 564.0, 774.2 and 998.4 non-conforming traces, respectively. We introduced

Table 4: Characteristics of the synthetic collection.

Per process model			Per event log		
Node Type	Avg.	Max.	Characteristic	Avg.	Max.
Places	19.1	77	Traces	1,000.0	1,000
Transitions	19.2	84	Trace length	6.4	122
And-splits	4.8	54	Uncertain traces	506.4	873
Xor-splits	3.4	28	Uncertain events	20.1%	68.7%
Silent steps	6.9	64	Resolutions	4.0	16,384

ordering uncertainty by abstracting all timestamps to minutes, i.e., by omitting all information on seconds and milliseconds. As a result, about 50.6% of the traces in the event logs become uncertain, having an average of 4.0 possible resolutions per trace, up to a maximum of 16,384.

The model-log pairs included in the real-world and synthetic collections differ considerably in terms of model complexity, trace length, degree of non-conformance, degree of uncertainty, and number of possible resolutions. As a result, these collections enable us to assess the impact of such key factors on the conformance checking accuracy and efficiency of our proposed behavioral models and approximation method. This gives the experimental results a high external validity.

6.2. Setup

To conduct our evaluation experiments, we implemented the proposed approach as a plug-in for the Java-based open-source Process Mining Framework *ProM 6*.[3]

Behavioral Models. Using the implementation and the datasets described above, we conducted experiments with the following behavioral models:

- *Trace equivalence:* The *trace equivalence model* from Section 4.1.
- *2-gram, 3-gram, 4-gram:* The *N-gram model* from Section 4.2. We use $N = 2$ (the strictest notion), $N = 3$, and $N = 4$.
- *Weak order:* The *weak order model* from Section 4.3.

Note that when a behavioral model returns a probability of zero for all possible resolutions of a trace, which happens if the model cannot derive any evidence from the event log, we regard all resolutions to be equally likely. That is, we assign a uniform probability $|\Phi(\sigma)|^{-1}$ to each resolution.

Approximation methods. To evaluate the impact of our proposed method for result approximation, we compute results based on three configurations:

---

[3]See http://www.promtools.org/

- *No approximation*: the conformance for all resolutions with a non-zero probability is assessed.

- *α=0.99 approximation*: the approximation method from Section 5 with $\alpha = 0.99$ and $\delta = 0.10$.

- *α=0.95 approximation*: the approximation method from Section 5 with $\alpha = 0.95$ and $\delta = 0.10$.

Performance measures. To determine how accurate the obtained conformance values are, we compare them to the true conformance values, i.e., a gold standard value, based on the order in which a trace's events appear in a log. To quantify the difference between *obtained* and *true* conformance values, we employ the Root-Mean Squared Error (RMSE). Given an event log L, a process model M, and a behavioral model B, the RMSE is given as follows for *probabilistic conformance*:

$$error(L, M, B) = 1 - \sqrt{\frac{\sum_{\sigma \in L}(conf(\sigma, M) - P^B_{conf}(\sigma, M))^2}{|L|}} \tag{11}$$

For the purposes of this evaluation, we consider conformance in terms of (weighted) fitness, as defined in Section 3.2.

Baseline. As a basis for comparison, we employ a baseline technique that considers each potential resolution of an uncertain trace to be equally likely, such as proposed by Lu et al. [22]. It assigns a *uniform probability* of $|\Phi(\sigma)|^{-1}$ to each resolution. The comparison against this baseline is intended to demonstrate the value of using probabilistic models that assess the likelihood of different resolutions.

6.3. Results

This section presents results on the accuracy of conformance checking using the proposed behavioral models (Section 6.3.1), the computational efficiency and accuracy of the approximation method (Section 6.3.2), and concludes with a discussion on when to select which model and method (Section 6.3.3).

6.3.1. Result Accuracy

Real-world collection. Figure 3 visualizes the conformance checking accuracy obtained by the behavioral models for the real-world event logs. The figure shows clear differences that exist across the various behavioral models, as well as across the three event logs. For the *BPI-12 log*, the weak order model performs best, achieving an average error of just 0.01. This accuracy is closely followed by the baseline (error of 0.02), whereas the other behavioral models perform considerably worse, ranging up to an error of 0.42 for the trace equivalence model. By contrast, the accuracy of all models for the *BPI-14 log* is comparable, ranging from 0.30 for the weak order and baseline models up to 0.38 for the 2-gram model. The *traffic fines log*,

again, shows considerable differences between the behavioral models. The baseline model achieves the worst accuracy (error of 0.17), while the trace equivalence and 2-gram models perform best (error of 0.08).
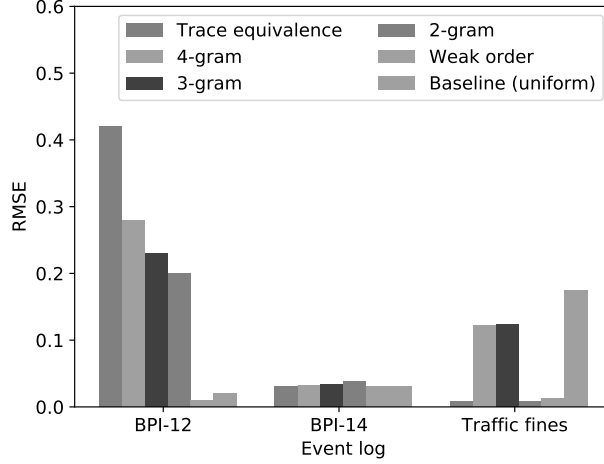

Figure 3: Result accuracy for real-world collection.

Synthetic collection. Figure 4 visualizes the results obtained for the 500 generated process models. The figure clearly shows that the relative performance of the behavioral models is, on average, stable across the four considered noise levels. The 2-gram model performs best, ranging from an error of 0.025 at the lowest noise level up to 0.36 for the highest level. When considering the accuracy of the behavioral models over all noise levels, we observe that the 2-gram model performs best, achieving an average RMSE of 0.031, followed by the 3-gram (0.034), weak order (0.038), 4-gram (0.041), and trace equivalence model (0.042). The baseline achieves an average RMSE of 0.078. This means that the 2-gram model reduces the average error by 60.2% compared to the baseline.

Considering the results obtained per event log, we observe a degree of variability among the performance of the behavioral models. The behavioral model that obtains the highest accuracy for an event log differs across the 2,000 model-log pairs. The 2-gram model is most often the best performing one (in 871 or 43.5% of the logs), respectively followed by the 3-gram (443 logs), trace equivalence (308), weak order (256), and the 4-gram model (122). Noticeably, there is no event log in the synthetic collection for which the baseline outperforms the proposed models.

6.3.2. Computational Efficiency

To assess the runtime efficiency of conformance checking under uncertainty and our proposed approximation method, we conducted experiments on a Linux server with an 8 Quad-Core-AMD 8384 processor, 2,7 GHz,
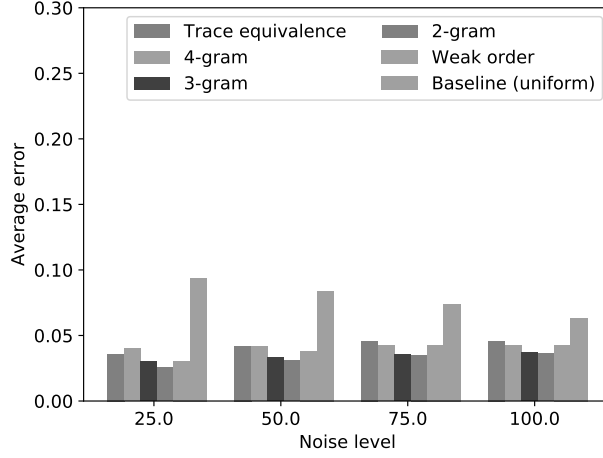
Figure 4: Result accuracy for synthetic collection.

using a 16GB Java Virtual Machine.

Real-world collection. First and foremost, we note that approximation has minimal impact on the conformance checking accuracy. For all real-world event logs, the error values obtained when using our approximation method are equal up to 3 decimal places compared to the results obtained without using approximation.

This near-identical accuracy is obtained while achieving considerable improvements in terms of runtime, as illustrated in Figure 5. As shown, the absolute runtimes clearly differ across the three event logs: Assessing the BPI-14 log takes several days, whereas the the traffic fines log takes a matter of seconds. The comparison between the original approach (No approx.) and the two configurations with approximation show notable runtime differences. In particular, for the BPI-12 and BPI-14 event logs, the approximation-based methods respectively save 89.0% and 90.1% of the runtime compared to the approach without approximation. For the traffic fines log, the difference is 30.1%.

Synthetic collection. Table 5 presents the runtime results obtained for the synthetic data collection.[4] We observe that the approximation method achieves considerable time savings, averaging between 58.9% and 45.9% savings per log and a maximum of 85.7% time saved. These gains are achieved while having a limited impact on the conformance checking accuracy of the approach. On average, the additional error incurred based on approximation is between only 0.02% and 0.03%, whereas the maximum additional error is 1.0%.

It is interesting to observe that the approximation method is actually applied to only about 2.0% of the traces with uncertainty, with a maximum of 18.8% in a single log. However, as shown by the much larger

---

[4]Due to near-identical performance of the configurations, we only report on the $\alpha = 0.01$ configuration.
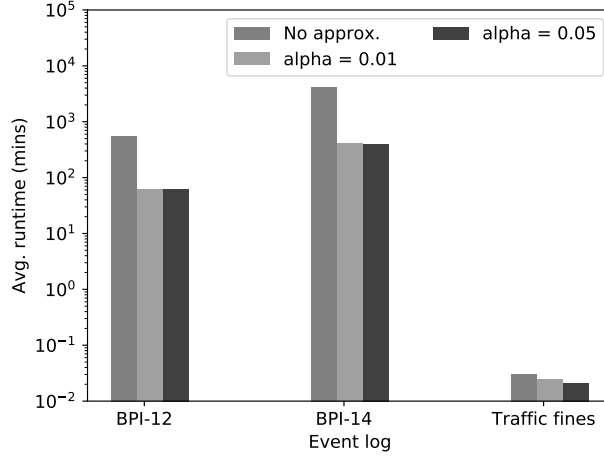
Figure 5: Runtime efficiency for the real-world collection.

gains in runtime, it is clear that the approximation method is applied to traces that contribute the most to the total runtime. This naturally follows from the approximation method's dependence on the central limit theorem, which enforces that approximation can only be applied to traces with at least 20 possible resolutions. In other words, the approximation method is generally applied to traces with the otherwise largest runtime.

Table 5: Runtime efficiency on the synthetic collection (averages over collection).

| Measure | Noise level | | | |
	25.0	50.0	75.0	100.0
Time, no approx. (avg.)	201.2m	212.6m	229.3m	232.9m
Time, with approx. (avg.)	82.6m	99.8m	112.0m	126.0m
Time saved (avg.)	58.9%	53.1%	51.2%	45.9%
Time saved (max.)	85.7%	80.9%	84.6%	80.7%
Traces approximated (avg.)	1.7%	1.9%	2.0%	2.2%
Traces approximated (max.)	17.2%	18.7%	17.9%	18.8%
Approx. error (avg.)	0.03%	0.02%	0.03%	0.02%
Approx. error (max.)	1.00%	0.98%	1.01%	1.00%

6.3.3. Behavioral Model Selection

The results presented for both the real-world and the synthetic data collections show that the performance of the various behavioral models differs per event log. A factor that plays an important role for the performance of the behavioral models is the amount of information that the model can utilize from the available event log. The higher the abstraction level of a behavioral model, the more information that can be taken into account. The trace equivalence model, the least abstract model, can only derive probabilistic

22

information from traces that do not contain any uncertainty. Even if such traces are available, the model can only provide probabilistic insights if the same trace is repeated for the process. This makes the trace equivalence model inapplicable to logs with a high variety and uncertainty, such as the BPI-14 log, in which over 97% of the traces have uncertainty. This problem can also occur for the 4-gram and 3-gram models, which still depend on repeated sequences of events without uncertainty.

Overall, we can conclude that the weak order model and 2-gram model are the best performing behavioral models. The weak order model performs best for 2 out of 3 real-world (BPI-12 and BPI-14) logs and 871 of the synthetic logs. The 2-gram model performs best for the other real-world log (traffic fines) and for 256 synthetic ones. Again, the performance of these models depends on their ability to derive information from a specific event log. Even though these models are the least restrictive in this regard, there are still traces for which the 2-gram model is unable to derive probabilistic information for any of its possible resolutions. This is, in particular, the case for the BPI-12 log, where the 2-gram model is only able to provide probabilistic information for 50.2% of the uncertain traces. As a result, the accuracy (error of 0.20) is considerably lower than that of the weak order model (error of 0.01), which is able to compute probabilistic information for 96.9% of the uncertain traces. Therefore, the evaluation results suggest to employ the 2-gram model for event logs where this model is able to derive sufficient information, whereas otherwise the weak order model represents the best alternative.

Aside from the selection of a behavioral model, the results from Section 6.3.2 show that the proposed approximation method can be used with little impact on the approach's accuracy, while gaining considerable benefits in terms of runtime. The preservation of such a high level of result accuracy is due to the method's grounding in statistical distributions, which requires that at least 20 possible resolutions of a trace are checked before applying approximation.

7. Related Work

In this section we discuss the three primary streams to which our work relates: conformance checking, sequence classification, and uncertain data management.

Conformance checking is commonly based on alignments, as discussed in Section 2.1. Yet, it may also be grounded in a comparison of sets of binary relations defined over events of a log and activities of a model, respectively [37]. Other work suggests to 'replay' traces in a process model, thereby identifying whether events denote valid activity executions [9, 38]. These approaches provide trade-offs in terms of their computational efficiency and the type of conformance insights that they provide. Uncertainty can manifest

23

itself in several ways in the context of conformance checking. Aside from the order uncertainty considered in this work, uncertainty can also follow from unknown mappings of events to process model activities [39, 40] and from the use of ambiguous process specifications [41].

Due to the computational complexity of alignment-based conformance checking, various angles have been followed to improve its runtime performance. Efficiency improvements have been obtained through the use of search-based methods [42, 43], planning algorithms [44], and distributed computing [45, 46]. Similar to the approximation method presented in Section 5, several approaches approximate conformance results to gain efficiency, e.g., by employing approximate alignments [47], sampling strategies [48], and applying divide-and-conquer schemes in the computation of conformance results [16, 49, 50]. These approaches can be regarded as complimentary to our approximation method, since our method reduces the number of resolutions for which conformance results need to be obtained, whereas these other approaches improve the efficiency achieved for individual resolutions.

Sequence classification refers to the task of assigning class labels to sequences of events [51], a task with applications, e.g., in genomic analysis, information retrieval, and anomaly detection. A key challenge of sequence classification is the high dimensionality of features, resulting from the sequential nature of events. A broad variety of sequence classification methods have been developed, which can be generally categorized as: (i) *feature-based* methods that transform sequences into feature vectors to apply conventional classification techniques [52]; (ii) *distance-based* methods that compute similarity among sequences [53]; and (iii) *model-based* methods that employ techniques such as Hidden Markov Models [54].

Some techniques for sequence classification resemble the behavioral models proposed in this work. Yet, there is a fundamental difference in the addressed problem: Sequence classification assigns labels to sequences, whereas our models aim at determining the actual sequence of events itself.

Data uncertainty is inherent to various application contexts, typically caused by data randomness, incompleteness, or limitations of measuring equipment [55]. This has created a need for algorithms and applications for uncertain data management [56]. Respective results include models for probabilistic and uncertain databases, see [57, 58], along with respective query mechanisms [59–61]. Moreover, models to capture uncertain event occurrence. This includes models where the event occurrence itself is uncertain [62] as well as those where uncertainty relates only to the time of event occurrence [63], which may be defined by a density function over an interval. Our model, in turn, incorporates the particular aspect of order uncertainty within a trace. The reason being that the probability of an event occurrence at a particular time is of minor importance when assessing conformance of a trace with respect to the execution sequences of a model.

8. Conclusion

In this paper, we overcame a key assumption of conformance checking that limits its applicability in practical situations: The requirement that the events observed for a case in a process are totally ordered. In particular, we established various behavioral models, each incorporating a different notion of behavioral abstraction. These behavioral models enable the resolution of partial orders by inferring probabilistic information from other process executions. Our evaluation of the approach based on real-world and synthetic data reveal that our approach achieves considerably more accurate results than an existing baseline, reducing the average error by 60.2%. To cope with the runtime complexity of conformance checking using uncertain event data, we also presented a sample-based approximation method. The conducted experiments demonstrate that this method can lead to runtime reductions of up to 90.1%, while still obtaining a near-identical conformance checking accuracy.

In our experiments, we focused on conformance on the trace level. However, as indicated in Section 3.2, the assignment of probabilities to resolutions is also beneficial for advanced feedback on non-conformance. Measures that quantify conformance may be computed per resolution and then be aggregated based on the assigned probabilities. Similarly, our probabilistic model highlights the importance of particular deviations identified on the level of resolutions. However, we see open research questions related to the perception of such probabilistic results in conformance checking and intend to explore this aspect in future case studies. We also aim to extend our work with behavioral models that go beyond temporal aspects of event logs. For instance, employing decision mining techniques, multi-perspective models may enable more fine-granular selection of traces for partial order resolution.

Reproducibility: Links to the source code and datasets required to reproduce the experiments are given in Section 6.

References

[1] M. Dumas, M. L. Rosa, J. Mendling, H. A. Reijers, Fundamentals of Business Process Management, Second Edition, Springer, 2018.

[2] W. M. P. V. der Aalst, Process Mining - Data Science in Action, Second Edition, Springer, 2016.

[3] W. Kettinger, J. Teng, S. Guha, Business Process Change: a Study of Methodologies, Techniques, and Tools, MIS quarterly (1997) 55–80.

[4] J. Carmona, B. F. van Dongen, A. Solti, M. Weidlich, Conformance Checking - Relating Processes and Models, Springer, 2018. `doi:10.1007/978-3-319-99414-7`.

[5] H. Schonenberg, R. Mans, N. Russell, N. Mulyar, W. Van der Aalst, Process flexibility: A survey of contemporary approaches, in: EOMAS, Springer, 2008, pp. 16–30.

[6] F. Bagayogo, A. Beaudry, L. Lapointe, Impacts of IT acceptance and resistance behaviors: a novel framework, in: Intl. Conf. Information Systems, 2013.

[7] R. Lu, S. Sadiq, G. Governatori, Compliance aware business process design, in: Intl. Conf. Business Process Management, Springer, 2007, pp. 120–131.

[8] F. Caron, J. Vanthienen, B. Baesens, Comprehensive rule-based compliance checking and risk management with process mining, Decision Support Systems 54 (3) (2013) 1357–1369.

[9] A. Rozinat, W. M. P. Van der Aalst, Conformance checking of processes based on monitoring real behavior, Inf. Syst. 33 (1) (2008) 64–95.

[10] W. Van der Aalst, K. Van Hee, J. M. Van der Werf, A. Kumar, M. Verdonk, Conceptual model for online auditing, Decision Support Systems 50 (3) (2011) 636–647.

[11] R. Accorsi, T. Stocker, On the exploitation of process mining for security audits: the conformance checking case, in: SAC, 2012, pp. 1709–1716.

[12] A. Adriansyah, B. van Dongen, W. M. P. Van der Aalst, Conformance checking using cost-based fitness analysis, in: EDOC, IEEE, 2011, pp. 55–64.

[13] W. Van der Aalst, A. Adriansyah, B. van Dongen, Replaying history on process models for conformance checking and performance analysis, WIDM 2 (2) (2012) 182–192.

[14] J. C. J. C. Bose, R. S. Mans, W. M. P. Van der Aalst, Wanna improve process mining results?, in: CIDM, IEEE, 2013, pp. 127–134. `doi:10.1109/CIDM.2013.6597227`.

[15] R. S. Mans, W. M. P. Van der Aalst, R. J. Vanwersch, A. J. Moleman, Process mining in healthcare: Data challenges when answering frequently posed questions, in: ProHealth, Springer, 2013, pp. 140–153.

[16] J. Munoz-Gama, J. Carmona, W. v. d. Aalst, Single-entry single-exit decomposed conformance checking, Inf. Syst. 46 (2014) 102–122.

[17] S. Suriadi, R. Andrews, A. H. ter Hofstede, M. T. Wynn, Event log imperfection patterns for process mining: Towards a systematic approach to cleaning event logs, Inf. Syst. 64 (2017) 132–150.

[18] S. Suriadi, C. Ouyang, W. M. van der Aalst, A. H. ter Hofstede, Event interval analysis: Why do processes take time?, Decision Support Systems 79 (2015) 77–98.

[19] C. Batini, M. Scannapieco, Data Quality: Concepts, Methodologies and Techniques, Data-Centric Systems and Applications, Springer, 2006. `doi:10.1007/3-540-33173-5`.

[20] E. Koskinen, J. Jannotti, Borderpatrol: isolating events for black-box tracing, in: J. S. Sventek, S. Hand (Eds.), Proc. 2008 EuroSys Conference, ACM, 2008, pp. 191–203. `doi:10.1145/1352592.1352613`.

[21] C. Mutschler, M. Philippsen, Reliable speculative processing of out-of-order event streams in generic publish/subscribe middlewares, in: Proc. 7th Intl. Conf. Distr. Event-based Syst., 2013, pp. 147–158.

[22] X. Lu, R. S. Mans, D. Fahland, W. M. P. Van der Aalst, Conformance checking in healthcare based on partially ordered

event data, in: ETFA, IEEE, 2014, pp. 1–8.

[23] T. T. L. Tran, C. A. Sutton, R. Cocci, Y. Nie, Y. Diao, P. J. Shenoy, Probabilistic inference over RFID streams in mobile environments, in: Proc. 25th Intl. Conf. Data Engineering ICDE, 2009, pp. 1096–1107.

[24] A. Senderovich, A. Rogge-Solti, A. Gal, J. Mendling, A. Mandelbaum, The ROAD from sensor data to process instances via interaction mining, in: CAISE, 2016, pp. 257–273.

[25] J. Carmona, B. van Dongen, A. Solti, M. Weidlich, Conformance Checking: Relating Processes and Models, Springer, 2018.

[26] S. Chen, J. L. Moore, D. Turnbull, T. Joachims, Playlist prediction via metric embedding, in: SIGKDD, ACM, 2012, pp. 714–722.

[27] M. Mohri, F. Pereira, M. Riley, Weighted finite-state transducers in speech recognition, Computer Speech & Language 16 (1) (2002) 69–88.

[28] M. Rosenblatt, A central limit theorem and a strong mixing condition, Proceedings of the National Academy of Sciences of the United States of America 42 (1) (1956) 43.

[29] A. Agresti, B. A. Coull, Approximate is better than "exact" for interval estimation of binomial proportions, The American Statistician 52 (2) (1998) 119–126.

[30] S. L. Lohr, Sampling: Design and Analysis: Design and Analysis, Chapman and Hall/CRC, 2019.

[31] S. J. Leemans, D. Fahland, W. M. P. Van der Aalst, Discovering block-structured process models from event logs-a constructive approach, in: International conference on applications and theory of Petri nets and concurrency, Springer, 2013, pp. 311–329.

[32] Van Dongen, B.F. (Boudewijn), Bpi challenge 2012 (2012). `doi:10.4121/UUID:3926DB30-F712-4394-AEBC-75976070E91F`.

[33] B. Van Dongen, BPI Challenge 2014 (2014). `doi:10.4121/uuid:c3e5d162-0cfd-4bb0-bd82-af5268819c35`.

[34] M. M. De Leoni, F. F. Mannhardt, Road traffic fine management process (2015). `doi:10.4121/UUID: 270FD440-1057-4FB9-89A9-B699B47990F5`.

[35] T. Jouck, B. Depaire, Generating artificial data for empirical analysis of control-flow discovery algorithms: A process tree and log generator, BISE (2018) 1–18.

[36] A. Rogge-Solti, M. Weske, Prediction of remaining service execution time using stochastic petri nets with arbitrary firing delays, in: ICSOC, Springer, 2013, pp. 389–403.

[37] M. Weidlich, A. Polyvyanyy, N. Desai, J. Mendling, M. Weske, Process compliance analysis based on behavioural profiles, Inf. Syst. 36 (7) (2011) 1009–1025.

[38] S. K. vanden Broucke, J. De Weerdt, J. Vanthienen, B. Baesens, Determining process model precision and generalization with weighted artificial negative events, IEEE Transactions on Knowledge and Data Engineering 26 (8) (2014) 1877–1889.

[39] T. Baier, J. Mendling, M. Weske, Bridging abstraction layers in process mining, Inf. Syst. 46 (2014) 123–139.

[40] H. Van der Aa, H. Leopold, H. A. Reijers, Efficient process conformance checking on the basis of uncertain event-to-activity mappings, TKDE.

[41] H. Van der Aa, H. Leopold, H. A. Reijers, Checking process compliance against natural language specifications using behavioral spaces, Inf. Syst. 78 (2018) 83–95.

[42] D. Reißner, R. Conforti, M. Dumas, M. L. Rosa, A. Armas-Cervantes, Scalable conformance checking of business processes, in: OTM to Meaningful Int. Syst., 2017, pp. 607–627. `doi:10.1007/978-3-319-69462-7\_38`.

[43] B. F. van Dongen, Efficiently computing alignments - using the extended marking equation, in: Business Process Management, 2018, pp. 197–214. `doi:10.1007/978-3-319-98648-7\_12`.

[44] M. de Leoni, A. Marrella, How planning techniques can help process mining: The conformance-checking case, in: Italian Symp. on Advanced Database Syst., 2017, p. 283.

[45] J. Evermann, Scalable process discovery using map-reduce, IEEE TSC 9 (3) (2016) 469–481. `doi:10.1109/TSC.2014.2367525`.

[46] C. Luo, F. He, C. Ghezzi, Inferring software behavioral models with MapReduce, Sci. Comput. Program. 145 (2017) 13–36. `doi:10.1016/j.scico.2017.04.004`.

[47] F. Taymouri, J. Carmona, A recursive paradigm for aligning observed behavior of large structured process models, in: Business Process Management, Springer, 2016, pp. 197–214. `doi:10.1007/978-3-319-45348-4\_12`.

[48] M. Bauer, H. van der Aa, M. Weidlich, Estimating process conformance by trace sampling and result approximation, in: Business Process Management, Springer, 2019, pp. 179–197.

[49] W. M. P. V. der Aalst, H. M. W. Verbeek, Process discovery and conformance checking using passages, Fundam. Inform. 131 (1) (2014) 103–138. `doi:10.3233/FI-2014-1006`.

[50] S. J. J. Leemans, D. Fahland, W. M. P. V. der Aalst, Scalable process discovery and conformance checking, Software and System Modeling 17 (2) (2018) 599–631. `doi:10.1007/s10270-016-0545-x`.

[51] Z. Xing, J. Pei, E. Keogh, A brief survey on sequence classification, SIGKDD Explor 12 (1) (2010) 40–48.

[52] G. Dong, J. Pei, Sequence data mining, Vol. 33, Springer Science & Business Media, 2007.

[53] E. Keogh, S. Kasetty, On the need for time series data mining benchmarks: a survey and empirical demonstration, Data Mining and knowledge discovery 7 (4) (2003) 349–371.

[54] P. K. Srivastava, D. K. Desai, S. Nandi, A. M. Lynn, HMM-mode–improved classification using profile hidden markov models by optimising the discrimination threshold and modifying emission probabilities with negative training sequences, BMC bioinf. 8 (1) (2007) 104.

[55] J. Pei, B. Jiang, X. Lin, Y. Yuan, Probabilistic skylines on uncertain data, in: Proc. 33rd Intl. Conf. Very large data bases, 2007, pp. 15–26.

[56] C. C. Aggarwal, P. S. Yu, A survey of uncertain data algorithms and applications, IEEE Trans. Knowl. Data Eng. 21 (5) (2009) 609–623.

[57] A. D. Sarma, O. Benjelloun, A. Halevy, J. Widom, Working models for uncertain data, in: 22nd Intl. Conf. Data Eng., IEEE, 2006, pp. 7–7.

[58] L. Peng, Y. Diao, Supporting data uncertainty in array databases, in: SIGMOD Intl. Conf. Mgt. of Data, ACM, 2015, pp. 545–560.

[59] K. Yi, F. Li, G. Kollios, D. Srivastava, Efficient processing of top-k queries in uncertain databases with x-relations, IEEE Trans. Knowl. Data Eng. 20 (12) (2008) 1669–1682.

[60] R. Murthy, R. Ikeda, J. Widom, Making aggregation work in uncertain and probabilistic databases, IEEE Trans. Knowl. Data Eng. 23 (8) (2011) 1261–1273.

[61] X. Lian, L. Chen, Probabilistic group nearest neighbor queries in uncertain databases, IEEE Transactions on Knowledge and Data Engineering 20 (6) (2008) 809–824.

[62] C. Ré, J. Letchner, M. Balazinska, D. Suciu, Event queries on correlated probabilistic streams, in: J. T. Wang (Ed.), Proc. SIGMOD Intl. Conf. Mgt. of Data, SIGMOD, ACM, 2008, pp. 715–728. `doi:10.1145/1376616.1376688`.

[63] H. Zhang, Y. Diao, N. Immerman, Recognizing patterns in streams with imprecise timestamps, Inf. Syst. 38 (8) (2013) 1187–1211. `doi:10.1016/j.is.2012.01.002`.